

Perform sentiment analysis on the amazon alexa reviews

In [1]:

```
import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
import nltk
nltk.download('punkt') #for word tokenization
nltk.download('stopwords') #for removing or getting list of stopwords
nltk.download('wordnet') #for Lemmatization
```

```
[nltk_data] Downloading package punkt to
[nltk_data]   C:\Users\swapn\AppData\Roaming\nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data]   C:\Users\swapn\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data]   C:\Users\swapn\AppData\Roaming\nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

Out[2]:

True

In [3]:

```
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
import matplotlib.pyplot as plt
from wordcloud import WordCloud
```

In [4]:

```
df = pd.read_csv("alex reviews.csv")
```

In [5]:

```
df
```

Out[5]:

Unnamed: 0		verified_reviews	feedback
0	0	Love my Echo!	1
1	1	Loved it!	1
2	2	Sometimes while playing a game, you can answer...	1
3	3	I have had a lot of fun with this thing. My 4 ...	1
4	4	Music	1
...
3145	3145	Perfect for kids, adults and everyone in betwe...	1
3146	3146	Listening to music, searching locations, check...	1
3147	3147	I do love these things, i have them running my...	1
3148	3148	Only complaint I have is that the sound qualit...	1
3149	3149	Good	1

3150 rows × 3 columns

In [6]:

```
df.drop(["Unnamed: 0"],axis=1,inplace=True)
```

In [7]:

```
df
```

Out[7]:

	verified_reviews	feedback
0	Love my Echo!	1
1	Loved it!	1
2	Sometimes while playing a game, you can answer...	1
3	I have had a lot of fun with this thing. My 4 ...	1
4	Music	1
...
3145	Perfect for kids, adults and everyone in betwe...	1
3146	Listening to music, searching locations, check...	1
3147	I do love these things, i have them running my...	1
3148	Only complaint I have is that the sound qualit...	1
3149	Good	1

3150 rows × 2 columns

In [8]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3150 entries, 0 to 3149
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   verified_reviews 3150 non-null   object
1   feedback         3150 non-null   int64
dtypes: int64(1), object(1)
memory usage: 49.3+ KB
```

```
wc = WordCloud(width=800, height=800, background_color="white", min_font_size=10)
wc.generate(''.join(df[df['feedback']==0]['verified_reviews']))

plt.figure(figsize=(6,6))
plt.imshow(wc)
plt.axis("off")
plt.show()
```



```

wc = WordCloud(width=800, height=800, background_color="white", min_font_size=10)
wc.generate(''.join(df[df['feedback']==1]['verified_reviews']))

plt.figure(figsize=(6,6))
plt.imshow(wc)
plt.axis("off")
plt.show()

```



```
stop = stopwords.words("english")
def clean_text(text):
    tokens = word_tokenize(text.lower())
    # Filter only alphabets
    word_tokens = [t for t in tokens if t.isalpha()]
    clean_tokens = [t for t in word_tokens if t not in stop]
    lemma = WordNetLemmatizer()
    lemma_tokens = [lemma.lemmatize(t) for t in clean_tokens]
    return " ".join(lemma_tokens)
```

```
df['verified_reviews'] = df['verified_reviews'].apply(clean_text)
```

In [13]:

```
df['verified_reviews'].head()
```

Out[13]:

```
0          love echo
1          loved
2  sometimes playing game answer question correct...
3  lot fun thing yr old learns dinosaur control l...
4          music
Name: verified_reviews, dtype: object
```

In [14]:

```
x= df['verified_reviews']
y= df['feedback']
```

In [15]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3)
```

In [16]:

```
sent_len = []
for t in df['verified_reviews']:
    sent_len.append(len(word_tokenize(t)))
df['sent_len'] = sent_len
df.head()
```

Out[16]:

	verified_reviews	feedback	sent_len
0	love echo	1	2
1	loved	1	1
2	sometimes playing game answer question correct...	1	17
3	lot fun thing yr old learns dinosaur control l...	1	18
4	music	1	1

In [17]:

```
max(sent_len)
```

Out[17]:

245

In [18]:

```
np.quantile(sent_len, 0.95)
```

Out[18]:

40.0

In [19]:

```
max_len = 40
```

In [20]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.preprocessing.text import Tokenizer
# Creates dictionary and every unique word is given number key
from tensorflow.keras.preprocessing import sequence
# To perform the padding of the documents with zero's to make the length of the
# document common
from tensorflow.keras.layers import (LSTM, Dropout, Embedding, SimpleRNN, GRU)
# All the index numbers are converted to vectors using Embedding
# SimpleRNN allows to implement the RNN architecture - activation function -tanh
# Dropout - manage overfitting of model
```

In [21]:

```
# Tokenization
tok = Tokenizer(char_level=False, split=" ")
tok.fit_on_texts(x_train)
```

In [22]:

```
tok.index_word
```

Out[22]:

```
{1: 'love',
 2: 'echo',
 3: 'great',
 4: 'alexa',
 5: 'music',
 6: 'work',
 7: 'like',
 8: 'use',
 9: 'sound',
10: 'device',
11: 'one',
12: 'dot',
13: 'easy',
14: 'speaker',
15: 'set',
16: 'good',
17: 'get',
18: 'thing'.
```

In [23]:

```
vocab_len = len(tok.index_word)
vocab_len
```

Out[23]:

2987

In [24]:

```
seq_train = tok.texts_to_sequences(x_train)
seq_train
```

```
[1, 118, 121, 208, 30],
[1784,
 458,
 171,
 1243,
 162,
 39,
 1785,
 119,
 575,
 837,
 14,
 105,
 114,
 535,
 37,
 9,
 213,
 16,
 126,
```

In [25]:

```
seq_padded_train = sequence.pad_sequences(seq_train, maxlen=max_len)
seq_padded_train
```

Out[25]:

```
array([[ 0,  0,  0, ...,  0,  1,  54],
       [ 0,  0,  0, ..., 22, 237, 1780],
       [ 0,  0,  0, ..., 26, 836, 1783],
       ...,
       [ 0,  0,  0, ...,  0,  0,  90],
       [ 0,  0,  0, ...,  0, 1639, 134],
       [ 0,  0,  0, ..., 26,  4, 107]])
```


In [26]:

```

model = Sequential()
# vectorization
model.add(Embedding(vocab_len+1,40, input_length=max_len, mask_zero=True))
# RNN Layer
model.add(SimpleRNN(32, activation="tanh"))
# ANN's hidden Layer
model.add(Dense(32, activation="relu"))
# To check on overfitting
model.add(Dropout(0.2))
# output Layer
model.add(Dense(1, activation="sigmoid"))

```

In [27]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 40, 40)	119520

simple_rnn (SimpleRNN)	(None, 32)	2336

dense (Dense)	(None, 32)	1056

dropout (Dropout)	(None, 32)	0

dense_1 (Dense)	(None, 1)	33
=====		
Total params: 122,945		
Trainable params: 122,945		
Non-trainable params: 0		

In [28]:

```
model.compile(loss="binary_crossentropy", optimizer="adam")
```

In [29]:

```
model.fit(seq_padded_train, y_train, batch_size=50, epochs=50)
```

```
Epoch 1/50
45/45 [=====] - 2s 11ms/step - loss: 0.5967
Epoch 2/50
45/45 [=====] - 0s 10ms/step - loss: 0.3192
Epoch 3/50
45/45 [=====] - 0s 10ms/step - loss: 0.2379
Epoch 4/50
45/45 [=====] - 0s 10ms/step - loss: 0.1135
Epoch 5/50
45/45 [=====] - 0s 10ms/step - loss: 0.0485
Epoch 6/50
45/45 [=====] - 0s 11ms/step - loss: 0.0324
Epoch 7/50
45/45 [=====] - 1s 12ms/step - loss: 0.0245
Epoch 8/50
45/45 [=====] - 1s 12ms/step - loss: 0.0177
Epoch 9/50
45/45 [=====] - 0s 9ms/step - loss: 0.0165
Epoch 10/50
45/45 [=====] - 0s 9ms/step - loss: 0.0182
Epoch 11/50
45/45 [=====] - 1s 11ms/step - loss: 0.0144
Epoch 12/50
45/45 [=====] - 0s 11ms/step - loss: 0.0164
Epoch 13/50
45/45 [=====] - 0s 10ms/step - loss: 0.0149
Epoch 14/50
45/45 [=====] - 0s 10ms/step - loss: 0.0128
Epoch 15/50
45/45 [=====] - 0s 10ms/step - loss: 0.0125
Epoch 16/50
45/45 [=====] - 0s 9ms/step - loss: 0.0124
Epoch 17/50
45/45 [=====] - 0s 9ms/step - loss: 0.0137
Epoch 18/50
45/45 [=====] - 0s 11ms/step - loss: 0.0149
Epoch 19/50
45/45 [=====] - 0s 9ms/step - loss: 0.0127
Epoch 20/50
45/45 [=====] - 0s 10ms/step - loss: 0.0089
Epoch 21/50
45/45 [=====] - 0s 10ms/step - loss: 0.0144
Epoch 22/50
45/45 [=====] - 0s 10ms/step - loss: 0.0148
Epoch 23/50
45/45 [=====] - 0s 10ms/step - loss: 0.0111
Epoch 24/50
45/45 [=====] - 0s 10ms/step - loss: 0.0124
Epoch 25/50
45/45 [=====] - 0s 10ms/step - loss: 0.0131
Epoch 26/50
45/45 [=====] - 0s 10ms/step - loss: 0.0119
Epoch 27/50
45/45 [=====] - 0s 10ms/step - loss: 0.0097
Epoch 28/50
45/45 [=====] - 0s 9ms/step - loss: 0.0136
```

```
Epoch 29/50
45/45 [=====] - 0s 9ms/step - loss: 0.0112
Epoch 30/50
45/45 [=====] - 0s 10ms/step - loss: 0.0106
Epoch 31/50
45/45 [=====] - 0s 10ms/step - loss: 0.0148
Epoch 32/50
45/45 [=====] - 0s 9ms/step - loss: 0.0114
Epoch 33/50
45/45 [=====] - 0s 9ms/step - loss: 0.0161
Epoch 34/50
45/45 [=====] - 0s 9ms/step - loss: 0.0091
Epoch 35/50
45/45 [=====] - 0s 9ms/step - loss: 0.0131
Epoch 36/50
45/45 [=====] - 0s 9ms/step - loss: 0.0122
Epoch 37/50
45/45 [=====] - 0s 9ms/step - loss: 0.0107
Epoch 38/50
45/45 [=====] - 0s 9ms/step - loss: 0.0128
Epoch 39/50
45/45 [=====] - 0s 9ms/step - loss: 0.0063
Epoch 40/50
45/45 [=====] - 0s 9ms/step - loss: 0.0134
Epoch 41/50
45/45 [=====] - 0s 9ms/step - loss: 0.0140
Epoch 42/50
45/45 [=====] - 0s 9ms/step - loss: 0.0116
Epoch 43/50
45/45 [=====] - 0s 9ms/step - loss: 0.0132
Epoch 44/50
45/45 [=====] - 0s 9ms/step - loss: 0.0128
Epoch 45/50
45/45 [=====] - 0s 10ms/step - loss: 0.0144
Epoch 46/50
45/45 [=====] - 0s 10ms/step - loss: 0.0114
Epoch 47/50
45/45 [=====] - 0s 9ms/step - loss: 0.0100
Epoch 48/50
45/45 [=====] - 0s 10ms/step - loss: 0.0118
Epoch 49/50
45/45 [=====] - 0s 10ms/step - loss: 0.0131
Epoch 50/50
45/45 [=====] - 0s 9ms/step - loss: 0.0103
```

Out[29]:

<tensorflow.python.keras.callbacks.History at 0x1bd9979a5b0>

In [30]:

```
seq_test = tok.texts_to_sequences(x_test)
seq_test
```

Out[30]:

```
[[129, 4, 26, 149, 4, 2, 219, 7, 92, 44, 26],
 [1, 2, 13, 615, 537, 48, 450, 8, 716, 21, 153, 677, 944, 4, 1074],
 [43, 79, 21, 17, 72, 64, 205, 258, 22, 1049, 101, 60, 87, 29],
 [28, 1, 2, 12, 17, 101, 2762, 171, 313, 21, 5, 87],
 [1111, 91],
 [531,
 29,
 70,
 402,
 59,
 46,
 70,
 714,
 275,
 1539,
 900,
 30,
 268.]
```

In [31]:

```
seq_padded_test = sequence.pad_sequences(seq_test, maxlen=max_len)
seq_padded_test
```

Out[31]:

```
array([[ 0,  0,  0, ..., 92, 44, 26],
       [ 0,  0,  0, ..., 944, 4, 1074],
       [ 0,  0,  0, ..., 60, 87, 29],
       ...,
       [ 0,  0,  0, ..., 500, 8, 1158],
       [ 0,  0,  0, ..., 0, 0, 152],
       [ 0,  0,  0, ..., 124, 169, 1096]])
```

In [32]:

```
y_hat = model.predict(seq_padded_test)
```

In [33]:

```
# y_hat contains probability
y_hat = np.where(y_hat>=0.5, 1, 0)
```

In [34]:

```
from sklearn.metrics import classification_report  
print(classification_report( y_test, y_hat))
```

	precision	recall	f1-score	support
0	0.44	0.29	0.35	83
1	0.93	0.96	0.95	862
accuracy			0.90	945
macro avg	0.69	0.63	0.65	945
weighted avg	0.89	0.90	0.90	945

In [35]:

```

model = Sequential()
# vectorization
model.add(Embedding(vocab_len+1,40, input_length=max_len, mask_zero=True))
# RNN Layer
# model.add(SimpleRNN(32, activation="tanh"))
model.add(LSTM(32, activation="tanh"))
# ANN's hidden layer
model.add(Dense(32, activation="relu"))
# To check on overfitting
model.add(Dropout(0.2))
# output layer
model.add(Dense(1, activation="sigmoid"))
model.compile(loss="binary_crossentropy", optimizer="adam")
model.fit(seq_padded_train, y_train, batch_size=50, epochs=50)

```

```

Epoch 1/50
45/45 [=====] - 6s 17ms/step - loss: 0.6283
Epoch 2/50
45/45 [=====] - 1s 17ms/step - loss: 0.3489
Epoch 3/50
45/45 [=====] - 1s 17ms/step - loss: 0.1870
Epoch 4/50
45/45 [=====] - 1s 17ms/step - loss: 0.1322
Epoch 5/50
45/45 [=====] - 1s 19ms/step - loss: 0.0805
Epoch 6/50
45/45 [=====] - 1s 16ms/step - loss: 0.0659
Epoch 7/50
45/45 [=====] - 1s 16ms/step - loss: 0.0291
Epoch 8/50
45/45 [=====] - 1s 16ms/step - loss: 0.0218: 0s -
loss: 0.021
Epoch 9/50
45/45 [=====] - 1s 16ms/step - loss: 0.0181
Epoch 10/50
45/45 [=====] - 1s 17ms/step - loss: 0.0180
Epoch 11/50
45/45 [=====] - 1s 22ms/step - loss: 0.0167
Epoch 12/50
45/45 [=====] - 1s 21ms/step - loss: 0.0161
Epoch 13/50
45/45 [=====] - 1s 20ms/step - loss: 0.0134
Epoch 14/50
45/45 [=====] - 1s 20ms/step - loss: 0.0129
Epoch 15/50
45/45 [=====] - 1s 20ms/step - loss: 0.0166
Epoch 16/50
45/45 [=====] - 1s 17ms/step - loss: 0.0136
Epoch 17/50
45/45 [=====] - 1s 17ms/step - loss: 0.0127
Epoch 18/50
45/45 [=====] - 1s 17ms/step - loss: 0.0136
Epoch 19/50
45/45 [=====] - 1s 20ms/step - loss: 0.0091
Epoch 20/50
45/45 [=====] - 1s 18ms/step - loss: 0.0107
Epoch 21/50
45/45 [=====] - 1s 20ms/step - loss: 0.0129

```

```
Epoch 22/50
45/45 [=====] - 1s 18ms/step - loss: 0.0121
Epoch 23/50
45/45 [=====] - 1s 19ms/step - loss: 0.0133
Epoch 24/50
45/45 [=====] - 1s 17ms/step - loss: 0.0154: 1s -
1
Epoch 25/50
45/45 [=====] - 1s 16ms/step - loss: 0.0088
Epoch 26/50
45/45 [=====] - 1s 16ms/step - loss: 0.0127
Epoch 27/50
45/45 [=====] - 1s 16ms/step - loss: 0.0173
Epoch 28/50
45/45 [=====] - 1s 16ms/step - loss: 0.0145
Epoch 29/50
45/45 [=====] - 1s 16ms/step - loss: 0.0140
Epoch 30/50
45/45 [=====] - 1s 16ms/step - loss: 0.0095TA: 0s
- loss
Epoch 31/50
45/45 [=====] - 1s 20ms/step - loss: 0.0138
Epoch 32/50
45/45 [=====] - 1s 18ms/step - loss: 0.0323
Epoch 33/50
45/45 [=====] - 1s 17ms/step - loss: 0.0206
Epoch 34/50
45/45 [=====] - 1s 16ms/step - loss: 0.0211
Epoch 35/50
45/45 [=====] - 1s 16ms/step - loss: 0.0124
Epoch 36/50
45/45 [=====] - 1s 16ms/step - loss: 0.0140
Epoch 37/50
45/45 [=====] - 1s 16ms/step - loss: 0.0138
Epoch 38/50
45/45 [=====] - 1s 16ms/step - loss: 0.0103
Epoch 39/50
45/45 [=====] - 1s 16ms/step - loss: 0.0104
Epoch 40/50
45/45 [=====] - 1s 16ms/step - loss: 0.0101: 0s -
10
Epoch 41/50
45/45 [=====] - 1s 16ms/step - loss: 0.0117
Epoch 42/50
45/45 [=====] - 1s 17ms/step - loss: 0.0109: 0s -
10s
Epoch 43/50
45/45 [=====] - 1s 17ms/step - loss: 0.0099
Epoch 44/50
45/45 [=====] - 1s 17ms/step - loss: 0.0130
Epoch 45/50
45/45 [=====] - 1s 17ms/step - loss: 0.0116
Epoch 46/50
45/45 [=====] - 1s 20ms/step - loss: 0.0084
Epoch 47/50
45/45 [=====] - 1s 17ms/step - loss: 0.0109
Epoch 48/50
45/45 [=====] - 1s 20ms/step - loss: 0.0112
Epoch 49/50
45/45 [=====] - 1s 17ms/step - loss: 0.0109
```

Epoch 50/50

45/45 [=====] - 1s 17ms/step - loss: 0.0093

Out[35]:

<tensorflow.python.keras.callbacks.History at 0x1bd9b24db20>

In [36]:

```
y_hat = model.predict(seq_padded_test)
y_hat = np.where(y_hat>=0.5, 1, 0)
print(classification_report(y_test, y_hat))
```

	precision	recall	f1-score	support
0	0.65	0.34	0.44	83
1	0.94	0.98	0.96	862
accuracy			0.93	945
macro avg	0.80	0.66	0.70	945
weighted avg	0.91	0.93	0.92	945

In [37]:

```
model = Sequential()
# vectorization
model.add(Embedding(vocab_len+1,40, input_length=max_len, mask_zero=True))
# RNN Layer
model.add(GRU(32, activation="tanh"))
# ANN's hidden layer
model.add(Dense(32, activation="relu"))
# To check on overfitting/
model.add(Dropout(0.2))
# output layer
model.add(Dense(1, activation="sigmoid"))
model.compile(loss="binary_crossentropy", optimizer="adam")
model.fit(seq_padded_train, y_train, batch_size=50, epochs=50)
```

```
Epoch 1/50
45/45 [=====] - 5s 15ms/step - loss: 0.6134
Epoch 2/50
45/45 [=====] - 1s 15ms/step - loss: 0.2745
Epoch 3/50
45/45 [=====] - 1s 15ms/step - loss: 0.1893
Epoch 4/50
45/45 [=====] - 1s 15ms/step - loss: 0.1218
Epoch 5/50
45/45 [=====] - 1s 15ms/step - loss: 0.0736
Epoch 6/50
45/45 [=====] - 1s 15ms/step - loss: 0.0387
Epoch 7/50
45/45 [=====] - 1s 15ms/step - loss: 0.0255
Epoch 8/50
45/45 [=====] - 1s 16ms/step - loss: 0.0216
Epoch 9/50
45/45 [=====] - 1s 15ms/step - loss: 0.0153
Epoch 10/50
45/45 [=====] - 1s 16ms/step - loss: 0.0174
Epoch 11/50
45/45 [=====] - 1s 15ms/step - loss: 0.0164
Epoch 12/50
45/45 [=====] - 1s 15ms/step - loss: 0.0156
Epoch 13/50
45/45 [=====] - 1s 18ms/step - loss: 0.0156
Epoch 14/50
45/45 [=====] - 1s 17ms/step - loss: 0.0119
Epoch 15/50
45/45 [=====] - 1s 17ms/step - loss: 0.0146
Epoch 16/50
45/45 [=====] - 1s 16ms/step - loss: 0.0137
Epoch 17/50
45/45 [=====] - 1s 17ms/step - loss: 0.0138
Epoch 18/50
45/45 [=====] - 1s 16ms/step - loss: 0.0131
Epoch 19/50
45/45 [=====] - 1s 15ms/step - loss: 0.0122
Epoch 20/50
45/45 [=====] - 1s 15ms/step - loss: 0.0135
Epoch 21/50
45/45 [=====] - 1s 17ms/step - loss: 0.0130
Epoch 22/50
45/45 [=====] - 1s 17ms/step - loss: 0.0143
```

```
Epoch 23/50
45/45 [=====] - 1s 17ms/step - loss: 0.0123
Epoch 24/50
45/45 [=====] - 1s 16ms/step - loss: 0.0159
Epoch 25/50
45/45 [=====] - 1s 16ms/step - loss: 0.0127
Epoch 26/50
45/45 [=====] - 1s 17ms/step - loss: 0.0155 ETA: 0s
- 1
Epoch 27/50
45/45 [=====] - 1s 17ms/step - loss: 0.0123
Epoch 28/50
45/45 [=====] - 1s 16ms/step - loss: 0.0129
Epoch 29/50
45/45 [=====] - 1s 15ms/step - loss: 0.0107
Epoch 30/50
45/45 [=====] - 1s 15ms/step - loss: 0.0130
Epoch 31/50
45/45 [=====] - 1s 19ms/step - loss: 0.0103
Epoch 32/50
45/45 [=====] - 1s 19ms/step - loss: 0.0115: 0s - 1
oss: 0.0
Epoch 33/50
45/45 [=====] - 1s 18ms/step - loss: 0.0143
Epoch 34/50
45/45 [=====] - 1s 20ms/step - loss: 0.0124
Epoch 35/50
45/45 [=====] - 1s 17ms/step - loss: 0.0120: 0s -
Epoch 36/50
45/45 [=====] - 1s 19ms/step - loss: 0.0164
Epoch 37/50
45/45 [=====] - 1s 16ms/step - loss: 0.0135
Epoch 38/50
45/45 [=====] - 1s 17ms/step - loss: 0.0135
Epoch 39/50
45/45 [=====] - 1s 17ms/step - loss: 0.0091
Epoch 40/50
45/45 [=====] - 1s 15ms/step - loss: 0.0182
Epoch 41/50
45/45 [=====] - 1s 19ms/step - loss: 0.0111
Epoch 42/50
45/45 [=====] - 1s 19ms/step - loss: 0.0137
Epoch 43/50
45/45 [=====] - 1s 19ms/step - loss: 0.0104
Epoch 44/50
45/45 [=====] - 1s 17ms/step - loss: 0.0104
Epoch 45/50
45/45 [=====] - 1s 17ms/step - loss: 0.0098
Epoch 46/50
45/45 [=====] - 1s 17ms/step - loss: 0.0097
Epoch 47/50
45/45 [=====] - 1s 16ms/step - loss: 0.0089
Epoch 48/50
45/45 [=====] - 1s 16ms/step - loss: 0.0123
Epoch 49/50
45/45 [=====] - 1s 19ms/step - loss: 0.0088
Epoch 50/50
45/45 [=====] - 1s 16ms/step - loss: 0.0108
```

Out[37]:



In [38]:

```
y_hat = model.predict(seq_padded_test)
y_hat = np.where(y_hat>=0.5, 1, 0)
print(classification_report(y_test, y_hat))
```

	precision	recall	f1-score	support
0	0.63	0.39	0.48	83
1	0.94	0.98	0.96	862
accuracy			0.93	945
macro avg	0.79	0.68	0.72	945
weighted avg	0.92	0.93	0.92	945

In []: